

# Basics

Guy J. Abel

# R Console Start Up

```
R version 3.4.3 (2017-11-30) -- "Kite-Eating Tree"
Copyright (C) 2017 The R Foundation for Statistical Computing
Platform: x86_64-w64-mingw32/x64 (64-bit)
```

R is free software and comes with ABSOLUTELY NO WARRANTY.  
You are welcome to redistribute it under certain conditions.  
Type `'license()'` or `'licence()'` for distribution details.

R is a collaborative project with many contributors.  
Type `'contributors()'` for more information and  
`'citation()'` on how to cite R or R packages in publications.

Type `'demo()'` for some demos, `'help()'` for on-line help, or `'help.start()'` for an HTML browser interface to help.

Type `'q()'` to quit R.

&gt;

# R Console Start Up

- Note the `>` at the bottom.
- Whenever you see this symbol, it means that R is not doing anything and just waiting for your input.
- It's called the prompt.

# Very Basics

- We can enter R commands at the prompt

```
> 1 + 2  
[1] 3
```

- When R is running it stores everything in the computer's active memory.
  - R knows 1 and 2, as well as the basic operator +.
  - It creates an object in the active memory containing the results of the computation.
  - R prints the content of the results object to the screen.
- R let's you know it has finished when displaying in the square brackets []
  - The square brackets [] tells you how many results were calculated.
  - In our case this object contains only one value, i.e. R has finished and printed the first and only value [1] of the computation.
- You can scroll through previous commands you've entered by using the up ("→") and down ("↓") keys on your keyboard.

# R Basics

- R may sometimes display a + (also known as the continuation line).
  - R will not report the [], as it is not displaying any content from an object stored in memory.
  - Entered an incomplete command
  - R is waiting for more input.

$$\begin{array}{c} > 1 + \\ + \end{array}$$

- If you are not sure what's going on hit Esc or Ctrl-C.
  - It will tell R to forget it and bring back the prompt.

# R Basics

- If you know what R wants, then you can complete the computation:

```
> 1 +  
+ 2  
[1] 3
```

- The continuation line also commonly occurs with
  - Unclosed parenthesis ( )
  - Unclosed quotes ""

```
> 7 / (1 + 3  
+ )  
[1] 1.75  
> "Guy  
+ "  
[1] "Guy\n"
```

- Remember if you are in a syntactic hole... Esc.

# Calculator

- R understands the following basic operators:
  - 1 + and - for addition and subtraction
  - 2 \* and / for multiplication and division
  - 3 ^ for exponents
- R observes standard rules of operator precedence.
- You can use brackets ( ) if you are not sure, e.g.

```
> 7 / (1 + 3)
[1] 1.75
```

is not the same as this:

```
> 7 / 1 + 3
[1] 10
```

# Character Strings

- R also allows you to type in strings of characters (letters, words, phrases)
- If you do not use quotation marks R will think you are asking for an object

```
> "Guy"  
[1] "Guy"  
> 'Guy'  
[1] "Guy"
```

```
> Guy  
Error: object 'Guy' not found
```

- We will go into depth on objects a little later.
- For now remember text must go in quotation.



# Sending Messages

- The `print()` function prints output in the R console.
  - Often do not need to use `print()` except when creating functions (more on creating functions later)

```
> 0.2
[1] 0.2
> print(0.2)
[1] 0.2
>
> "Guy"
[1] "Guy"
> print("Guy")
[1] "Guy"
```

- The `paste()` function can link together character strings.
  - Also allows non-characters as inputs, converts output to one character string.

```
> paste("I ate", "a bar of chocolate")
[1] "I ate a bar of chocolate"
> paste("I ate", 1, "bar of chocolate")
[1] "I ate 1 bar of chocolate"
> paste("I ate", 1, "bar of chocolate", sep = " . ")
[1] "I ate . 1 . bar of chocolate"
```

# Comments

- The comment operator # will tell R to ignore everything printed after it (in the current line).
- Extremely useful to annotate your code.

```
> 1 + 2 + 3 # Here R does some serious Maths  
[1] 6
```

- It is good practice to annotate your code.
  - Within a surprisingly short period of time you will forget what each bit of code is trying to do.

# Comments

- Be careful.

```
> 1 + 2 # + 20  
[1] 3
```

- Misplaced comments can break your code.

```
> 1 + # 2  
+
```

# Spacing

- For the most part, R does not care about spacing.

```
> 1 + 2
[1] 3
>
> # same as
> 1      +          2
[1] 3
```

- For character strings spaces matter...

```
> print("Strings obey spacing.")
[1] "Strings obey spacing."
> print(" Strings obey          spacing . ")
[1] " Strings obey          spacing . "
```

# Semi-Colons

- R evaluates code line by line.
- A line break tells R that a statement is to be evaluated.
- Instead of a line break, you can use a semicolon (;) to tell R where statements end.

```
> "Guy"  
[1] "Guy"
```

```
> "Guy" "Abel"  
Error: unexpected string constant in ""Guy" "Abel""
```

```
> "Guy"; "Abel"  
[1] "Guy"  
[1] "Abel"
```

# Exercise 1 (ex11.R)

- Open ex11.R and complete the following exercises:

*# 1. 5 plus 6*

*# 2. 2 multiplied by 8*

*# 3. 8 divided by 3*

*# 4. 909 minus 506*

*# 5. 5 to the power of 10*

*# 6. Tell R to say your first name*

*# 7. Q: Which symbol does R use to ignore all code after?*

*# A:*

*# 8. On one line of code with two R print commands write  
# your first and last name.*

# Basic Functions

- R comes with a many many pre-installed functions.
  - Installed as part of the base package which is located in your library directory.
  - Functions have names and take arguments in parentheses: `function(...)`
- Takes the form

```
function(argument1, argument2)
```

- Arguments are options for the function.
- Each function has different arguments.
- If there are multiple arguments, use a comma , to separate.

# Basic Functions

- We can find out what arguments (options) are for a function using ?

```
> ?log
```

- In RStudio you can also use Tab once you have typed the function name and opened the parenthesis
- The help file reports that the function takes two arguments, separated by a comma.
  - `x` a number you want to take the logarithm of (no default)
  - `base` the base system for the logarithm (default base = `exp(1)`)

```
> log(x = 10)
[1] 2.302585
> # same as
> log(x = 10, base = exp(1))
[1] 2.302585
> # change the base argument...
> log(x = 10, base = 10)
[1] 1
```



# Basic Functions

- R knows what arguments are supplied to the function and their order.

```
> log(x = 10)
[1] 2.302585
> # same as
> log(10)
[1] 2.302585
>
> # knows the order as well (i.e. second input is the `base`)
> log(10, 10)
[1] 1
> # same as
> log(x = 10, base = 10)
[1] 1
```

- Whilst this reduces your typing, I advise not to do it.
  - Difficult to remember which inputs for which arguments.
  - Even more difficult for others looking at your code.
  - Using Tab and Enter in RStudio means its easy to write argument names fully.

# Basic Functions

## Basic mathematical functions in base R

Function	Description
<code>log()</code>	computes natural logarithms
<code>log10()</code>	computes logarithm to the base 10
<code>exp()</code>	computes the exponential function
<code>sqrt()</code>	takes the square root
<code>abs()</code>	returns the absolute value
<code>sin()</code>	returns the sine
<code>factorial()</code>	returns the factorial
<code>sign()</code>	returns the sign (negative or positive)
<code>round()</code>	rounds the input to the desired digit

Remember we can use `?` to understand how any function works, e.g. `?round`

## Exercise 2 (ex12.R)

*# 1. Natural logarithm of 5*

*# 2. Square root of 121*

*# 3. Absolute value of 10 and -11*

*# 4. 8 x 7 x 6 x 5 x 4 x 3 x 2 x 1*

*# 5. Round pi (3.141593) to 3 decimal places*

*# 6. The sin angle of 60*

*# 7. e to the power of 4*

*# 8. Print your name*

# Logic Operators

- Among the most used features of R are logical operators.
- You will use these throughout your code and they are crucial for all sorts of data manipulation.
- When R evaluates statements containing logical operators it will return either TRUE or FALSE

Function	Description
<	less than
<=	less than or equal to
>	greater than
>=	greater than or equal to
==	equal
!=	not equal
&	and
	or

## Exercise 3 (ex13.R)

*# 1. Test if 3 is larger than pi*

*# 2. Test if pi is equal to 3.141593*

*# 3. Test if 7 divided by 3 is less than or equal to 3*

*# 4. Test if 5 times 2 is greater than or equal to 10*

*# 5. Test if 1 plus 5 is not equal to 7*

*# 6. Test if logarithm of 1000 is larger 7*

*# 7. Test if pi is greater than 3 and 4*

*# 8. Test if pi is less than 3 or 4*

# Assignment and Reference

- Instead of recalculating everything over and over again we can give things names and recall them later.
- Before we implicitly relied on and then manipulated objects and R implicitly printed these objects to the screen.

```
> 1 + 2  
[1] 3
```

- We can assign and recall names using the assignment operator `<-`.
  - Think of this as the M+ button on your calculator.

```
> a <- 1 + 2
```

- Above, R no longer gave the answer to our problem. It just returns the prompt.

# Assignment and Reference

- We can print the results to the screen by typing the name of our new object

```
> a  
[1] 3
```

- If you give R the name of some object it knows you don't even have to use the `print()` function.
- Just type in the name and R will give you the result.

```
> a  
[1] 3
```

- It does not do the calculation again when you print the object.

# Beware

- We can use just about any name we like except
  - Cannot be a number, e.g. `3 <- 1 + 2`
  - Cannot start with with a number, e.g. `3a <- 1 + 2`
- You can break R code by creating objects that already are existing functions.

```
> exp
function (x)  .Primitive("exp")
> exp <- 1 + 3
> exp
[1] 4
```

- The `exp` function still works, but is lower down in the search environment than the new object.
  - We could create our own function `exp` that does something completely different.
  - Directly get a function in a specific package using `::`

```
> base::exp
function (x)  .Primitive("exp")
```



# Beware

- R is case-sensitive.

```
> A
Error: object 'A' not found
```

```
> a
[1] 3
```

- If you are not sure, check the potential object name by printing it before assigning it!

```
> e
Error in eval(expr, envir, enclos): object 'e' not found
```

# Playing with Trivial Vectors

- Named objects behave just like the ones R already knows.
- This is very useful when things get more complicated:

```
> a
[1] 3
> a * 2
[1] 6
>
> b <- a * sqrt(a)
> b
[1] 5.196152
```

# Keeping Track of Objects

- To see what objects you have created (the ones R stored in active memory) you can use the `ls()` function.

```
> ls()
[1] "a"    "b"    "exp"
```

- RStudio also shows each created object in the Environment tab.
- If you want to remove an object from memory use the `rm()` function.
- Be very careful. This will delete thing permanently.

```
> rm(b, exp)
> ls()
[1] "a"
```

- If you want to remove all objects from active memory can also use the button with a brush on it above the object list in RStudio or directly in the console:

```
> rm(list = ls())
```

# Real Vectors

- So far we have only created trivial vectors of length 1. Let's assign some longer ones.
- To do this you will use the `c()` function.
  - The `c` stands for concatenate,
  - Combines a bunch of elements together, separated by commas.
  - Make sure you never call an object `c` (like we just created objects `a` and `b`)

```
> v1 <- c(1,2,3,4,5,6,7,8,9,10)
```

How about a character vector?

```
> v2 <- c("a", "b", "c", "d")
> v2
[1] "a" "b" "c" "d"
```

Or ...

```
> # will convert everything to characters if there is at least one
> v3 <- c(1, "two", 3, 4)
> v3
[1] "1" "two" "3" "4"
```

# Real Vectors

- You can also string multiple vectors together with the `c()` function.

```
> v4 <- c(v1 , v2 , v1)
```

```
> v4  
[1] "1" "2" "3" "4" "5" "6" "7" "8" "9" "10" "a" "b" "c" "d"  
[15] "1" "2" "3" "4" "5" "6" "7" "8" "9" "10"
```

# Simplifying Vector Creation

- Using the `c()` function can be tedious
  - Have to manually type all elements of a vector.
  - Will make mistakes
- Use the colon `(:)` to tell R to create an integer vector.

```
> 1:20
[1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20
> # backwards
> 10:-5
[1] 10 9 8 7 6 5 4 3 2 1 0 -1 -2 -3 -4 -5
```

- Use the `seq()` function for more general sequences.

```
> seq(from = 0, to = 10)
[1] 0 1 2 3 4 5 6 7 8 9 10
> # the 'by' argument let's you set the increments
> seq(from = 0, to = 10, by = 2)
[1] 0 2 4 6 8 10
```

# Simplifying Vector Creation

```
> # the length.out argument specifies the
> # length of the vector and R figures out
> # the increments itself
> seq(0, 10, length.out = 25)
 [1] 0.0000000 0.4166667 0.8333333 1.2500000 1.6666667 2.0833333
 [7] 2.5000000 2.9166667 3.3333333 3.7500000 4.1666667 4.5833333
[13] 5.0000000 5.4166667 5.8333333 6.2500000 6.6666667 7.0833333
[19] 7.5000000 7.9166667 8.3333333 8.7500000 9.1666667 9.5833333
[25] 10.0000000
```

# Simplifying Vector Creation

- Another useful function is `rep()` which allows you to repeat things.

```
> rep(x = 0, times = 10)
[1] 0 0 0 0 0 0 0 0 0 0
> # as always you can drop the argument name
> rep(x = "Hello", times = 3)
[1] "Hello" "Hello" "Hello"
> # repeating vec 1 twice
> rep(x = v1 , times = 2)
[1] 1 2 3 4 5 6 7 8 9 10 1 2 3 4 5 6 7 8 9 10
> # we can repeat each element as well
> rep(x = v2 , each = 2)
[1] "a" "a" "b" "b" "c" "c" "d" "d"
```



# Vector Operations

- R is very powerful when working with vectors.
- Most standard mathematical functions work with vectors.

```
> v1 + v1
[1] 2 4 6 8 10 12 14 16 18 20
> v1 ^ 2
[1] 1 4 9 16 25 36 49 64 81 100
> log(v1)
[1] 0.0000000 0.6931472 1.0986123 1.3862944 1.6094379 1.7917595 1.9459101
[8] 2.0794415 2.1972246 2.3025851
```

## Exercise 4 (ex14.R)

# 1. Create a vector called *a1* for the 1, 4, 9, 16, 25

# 2. Create a vector called *a2* from 1 to 100

# 3. Create a vector called *a3* from -8 to 20 in steps of 4

# 4. Create a vector called *a4* that combines *a1* and *a3*

# 5. Create a vector called *a5* that has your first and last name repeated 3 times

# 6. Find the square root of each element in *a1*

# 7. Which elements in *a3* are positive (greater than zero)

# 8. Divide *a2* by *a1*. Then divide *a2* by *a3*. What is going on?

# Selected Functions for Vectors

- R has many functions in the base package that work with vectors.

Function	Description
<code>sum()</code>	sums of the elements of the vector
<code>prod()</code>	product of the elements of the vector
<code>min()</code>	minimum of the elements of the vector
<code>max()</code>	maximum of the elements of the vector
<code>range()</code>	the range of the vector
<code>sort()</code>	sorts the vector (argument: <code>decreasing = FALSE</code> )
<code>rev()</code>	Reverse the order of the vector
<code>length()</code>	returns the length of the vector
<code>which()</code>	returns the index after evaluating a logical statement
<code>unique()</code>	returns a vector of all the unique elements of the input
<code>table()</code>	returns tabulation count for each unique element

## Exercise 5 (ex15.R)

```
# 1. Find the sum of all elements in a1

# 2. What is the largest element in a4

# 3. Multiply all elements of a2 together

# 4. How many elements are there in a5

# 5. Reverse the order of a2

# 6. Which elements of a1 are greater than 10

# 7. What is the range of values in a4

# 8. What are the unique elements of a4

# 9. Sort the values in a4

# 10. Create a tabulation of the elements in a4
```

# Selected Statistical Functions for Vectors

- R has many statistical functions in the `stats` package that work with vectors.

Function	Description
<code>mean()</code>	mean of the elements
<code>median()</code>	median of the elements
<code>sd()</code>	the standard deviation
<code>var()</code>	the variance (on
<code>cov()</code>	the covariance (takes two inputs <code>cov(x,y)</code> )
<code>cor()</code>	the correlation coefficient (takes two inputs <code>cor(x,y)</code> )
<code>IQR()</code>	Inter Quartile Range
<code>quantile()</code>	Sample quantiles corresponding to the given probabilities
<code>summary()</code>	returns summary statistics

# Sampling

- R can generate many different types of random numbers.
- The `sample()` function draws randomly from a given vector

```
> v1
[1] 1 2 3 4 5 6 7 8 9 10
> sample(x = v1, size = 5)
[1] 10 1 3 2 8
> sample(v1, size = 15, replace = TRUE)
[1] 8 10 4 5 10 4 9 6 8 7 5 7 2 5 5
```

Function	Description
<code>rnorm()</code>	random generation for the normal distribution (mean, sd)
<code>rbinom()</code>	random generation for the binomial distribution (size, prob)
<code>rpois()</code>	random generation for the Poisson distribution (lambda)
<code>runif()</code>	random generation for the uniform distribution (min, max)

## Exercise 6 (ex16.R)

*# 1. Create a vector a6 that is a sequence from 1 to 25 of length 8*

*# 2. What is the mean of a6*

*# 3. what is the standard deviation of a6*

*# 4. Create a vector a7 that is a random sample of size 8 from a2*

*# 5. What is the correlation of a6 and a7*

*# 6. Create a vector a8 of length 8 of random numbers from a normal  
distribution with mean 10 and standard deviation 2*

*# 7. What is Inter Quartile Range of a8*

*# 8. Create a summary of a8*

# Objects

- R is an object oriented language.
  - Everything in R is an object.
  - When R does anything, it creates and manipulates objects.
  - Objects are a bit like the memory button on the calculator, but with much much more freedom.
- R objects come in different types and flavors.



# Basic Objects: Vectors

- One-dimensional sequences of elements of the same mode. (More on modes later)
- For example, this could be vector of length 26 (i.e. one containing 26 elements) where each element is a letter in the alphabet.
- R has some built in vectors

```
> pi
[1] 3.141593
> LETTERS
[1] "A" "B" "C" "D" "E" "F" "G" "H" "I" "J" "K" "L" "M" "N" "O" "P" "Q"
[18] "R" "S" "T" "U" "V" "W" "X" "Y" "Z"
```

# Basic Objects: Matrices

- Two dimensional rectangular objects (matrices) or
- All elements of matrices have to be of the same mode.
- Can build matrices using the `matrix()` function

```
> matrix(data = 1:12, nrow = 4)
```

```
  [,1] [,2] [,3]
```

```
[1,]    1     5     9
```

```
[2,]    2     6    10
```

```
[3,]    3     7    11
```

```
[4,]    4     8    12
```

# Basic Objects: Lists

- Like vectors but they do not have to contain elements of the same mode.
- The first element of a list could be a vector of the 26 letters of the alphabet.
- The second element could contain a vector of numbers.
- A third could be a 2 by 3 matrix.

```
> list(LETTERS, 1:10, matrix(1:6, nrow = 2))
[[1]]
[1] "A" "B" "C" "D" "E" "F" "G" "H" "I" "J" "K" "L" "M" "N" "O" "P" "Q"
[18] "R" "S" "T" "U" "V" "W" "X" "Y" "Z"

[[2]]
[1] 1 2 3 4 5 6 7 8 9 10

[[3]]
      [,1] [,2] [,3]
[1,]    1    3    5
[2,]    2    4    6
```

# Basic Objects: Data Frames

- Data frames are best understood as special matrices (technically they are a type of list).
- They are two dimensional containers with
  - Rows corresponding to observations
  - Columns corresponding to vectors
- R has some built in data frames

```
> swiss
```

	Fertility	Agriculture	Examination	Education	Catholic
Courtelary	80.2	17.0	15	12	9.96
Delemont	83.1	45.1	6	9	84.84
Franches-Mnt	92.5	39.7	5	5	93.40
Moutier	85.8	36.5	12	7	33.77
Neuveville	76.9	43.5	17	15	5.16
Porrentruy	76.1	35.3	9	7	90.57
Broye	83.8	70.2	16	7	92.85
Glane	92.4	67.8	14	8	97.16
Gruyere	82.4	53.3	12	7	97.67
Sarine	82.9	45.2	16	13	91.38
Veveyse	87.1	64.5	14	6	98.61
Aigle	64.1	62.0	21	12	8.52
Aubonne	66.9	67.5	14	7	2.27
Avenches	68.9	60.7	19	12	4.43

# Modes

- All objects have a certain mode.
- Some objects can only deal with one mode at a time, others can store elements of multiple modes.
- Some basic modes include:
  - 1 integer: integers (e.g. 1, 2 or -1000)
  - 2 numeric: real numbers (e.g. 2.336, -0.35)
  - 3 character: elements made up of text-strings (e.g. "text", "Hello World!", or "123")
  - 4 logical: data containing logical constants (i.e. TRUE and FALSE)

# Indexing

- Sometimes you do not want to print or manipulate an entire vector.
- This is where indexing comes in.
- Indexing vectors is done with `[]`.

```
> v4 <- c("I", "really", "like", "chocolate", "ice cream.")
> # with the bracket we reference the third element
> v4[3]
[1] "like"
> # we can reference a sequence of elements
> v4[2:4]
[1] "really"      "like"        "chocolate"
> # or any elements we like
> v4[c(1,3,4)]
[1] "I"           "like"        "chocolate"
```

# Indexing

```
> # all except the 2nd element
> v4[-2]
[1] "I"           "like"         "chocolate"   "ice cream."
> # and we can change elements
> v4[4] <- "strawberry"
> v4
[1] "I"           "really"       "like"         "strawberry"  "ice cream."
```

# Indexing

- Indexing also works with matrices
- Using a , to separate rows and columns [row, column]

```
> # create a matrix
> m0 <- matrix(data = 1:12, nrow = 3)
> m0
      [,1] [,2] [,3] [,4]
[1,]    1    4    7   10
[2,]    2    5    8   11
[3,]    3    6    9   12
> # bottom right element
> m0[3, 4]
[1] 12
> # third column
> m0[, 3]
[1] 7 8 9
```



# Indexing

- The [row, column] indexing also works with data frames
- R users usually prefer to use the name of the column
- The dataframe\$column selects the entire column

```
> # can use the raw row and column coordinates
> swiss[38,1]
[1] 79.3
> # the row and column names
> swiss["Sion","Fertility"]
[1] 79.3
> # the column name and row coordinates
> # this is useful if there are no row names, which can often be the case.
> swiss$Fertility[38]
[1] 79.3
> # the whole column
> swiss$Fertility
[1] 80.2 83.1 92.5 85.8 76.9 76.1 83.8 92.4 82.4 82.9 87.1 64.1 66.9 68.9
[15] 61.7 68.3 71.7 55.7 54.3 65.1 65.5 65.0 56.6 57.4 72.5 74.2 72.0 60.5
[29] 58.3 65.4 75.5 69.3 77.3 70.5 79.4 65.0 92.2 79.3 70.4 65.7 72.7 64.4
[43] 77.6 67.6 35.0 44.7 42.8
```

# Missing Values

- NA is used to represent missing data
- Different functions treat missing data differently

```
> v5 <- c(10, 7, NA, NA, 0, NA, -2, 8)
> v5
[1] 10  7 NA NA  0 NA -2  8
> # this won't work by default
> max(v5)
[1] NA
> # there is usually an option to allow functions to omit NA
> max(v5, na.rm = TRUE)
[1] 10
> # test for missing values using is.na
> is.na(v5)
[1] FALSE FALSE  TRUE  TRUE FALSE  TRUE FALSE FALSE
> # test for non missing values using !is.na
> !is.na(v5)
[1]  TRUE  TRUE FALSE FALSE  TRUE FALSE  TRUE  TRUE
> # this works (does not have na.rm argument)
> summary(v5)
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   NA's
  -2.0     0.0     7.0     4.6     8.0    10.0     3
```

## Exercise 7 (ex17.R)

```
# 1. Print out the second row of the m1 matrix, where
# m1 <- matrix(1:28, nrow = 4, ncol = 7)

# 2. Print out the top right value of the m1 matrix

# 3. Print out the first column of the m1 matrix

# 4. Print out the whole of the Infant.Mortality column in the swiss data

# 5. What is the maximum fertility in the swiss data

# 6. Print out the Catholic value for Glane in the swiss data

# 7. Set the third and sixth element of a8 to missing values

# 8. What is the variance of the remaining numbers in a8
```

# Conditionals

- R is programmable language.
- Conditionals are one of the fundamentals of basic programming
  - Evaluating conditional statements (returns TRUE or FALSE)
  - Performing some action based on the evaluation.
- The `if()` function is used when you want action only when a statement is TRUE.
- Takes format

```
if(statement){  
  action  
}
```

- The statement must one logical value (TRUE or FALSE)

# Conditionals

```
> v6 <- rnorm(n = 10)
> v6
[1] -0.6264538 0.1836433 -0.8356286 1.5952808 0.3295078 -0.8204684
[7] 0.4874291 0.7383247 0.5757814 -0.3053884
>
> # test to see if the 2nd element is positive
> if(v6[2] > 0){
+   print("My number is positive")
+ }
[1] "My number is positive"
>
> # test to see if the 3rd element is positive
> if(v6[3] > 0){
+   print("My number is positive")
+ }
>
> # paste() function to join together numbers and character stings
> if(v6[2] > 0){
+   print(paste(v6[2], "is positive"))
+ }
[1] "0.183643324222082 is positive"
```

# Conditionals

- The `ifelse()` function is used when you want one action when statement TRUE and different action when FALSE.
- Very useful when creating variables in data frames. Has three arguments
  - test the statement
  - yes action if statement is TRUE
  - no action if statement is FALSE

```
> # logical test
> v6 > 0
[1] FALSE TRUE FALSE TRUE TRUE FALSE TRUE TRUE TRUE FALSE
>
> # display a character vector for results
> ifelse(test = v6 > 0, yes = "Positive", no = "Negative")
[1] "Negative" "Positive" "Negative" "Positive" "Positive" "Negative"
[7] "Positive" "Positive" "Positive" "Negative"
>
> # can return numeric values as well
> ifelse(test = v6 > 0, yes = 0, no = 1)
[1] 1 0 1 0 0 1 0 0 0 1
```

# Conditionals

```
> # remember v5
> v5
[1] 10  7 NA NA  0 NA -2  8
>
> # remember is.na()
> is.na(v5)
[1] FALSE FALSE  TRUE  TRUE FALSE  TRUE FALSE FALSE
>
> # can use any test that returns logical
> ifelse(test = is.na(v5), yes = "Missing", no = "Observed")
[1] "Observed" "Observed" "Missing"  "Missing"  "Observed" "Missing"
[7] "Observed" "Observed"
```

# Loops

- Loops are functions that carry out repetitive actions.
- The `for()` function in R is used to repeat actions a given number of times
- Takes format:

```
for(element in vector){  
  action  
}
```

- Within the `{}` we put the actions we want repeated multiple times.
  - General code, using `element` in place of the parts of the code that are non-general.
- Within the `()` we put an instruction.
  - The `element` values are taken from `vector` in each round of the loop.
  - Typically people use a single letter such as `i`.
  - The `vector` contains all the possible `element` values to be used.



# Loops

- For each  $i$  in the vector  $v7$  we calculate  $x$  and ask R to print to the screen.
- Or put differently, for each  $i$  element in a vector we do the thing inside the curly braces.

```
> v7 <- rnorm(n = 5)
> v7
[1] 1.5117812 0.3898432 -0.6212406 -2.2146999 1.1249309
> for(i in v7){
+   x <- i * 2
+   print(x)
+ }
```

```
[1] 3.023562
[1] 0.7796865
[1] -1.242481
[1] -4.4294
[1] 2.249862
```

# Loops

```
> v7
[1] 1.5117812 0.3898432 -0.6212406 -2.2146999 1.1249309
> # basic action with no print
> for(i in v7){
+   x <- i * 2
+ }
> # still does the work just not printed.
> # only saves the last version of x
> x
[1] 2.249862
```

# Loops

```
> # create empty vector to save action
> n <- length(v7)
> n
[1] 5
> v8 <- rep(NA, times = n)
> v8
[1] NA NA NA NA NA
> for(i in 1:n){
+   v8[i] <- v7[i] * 2
+ }
> v8
[1] 3.0235623 0.7796865 -1.2424812 -4.4293998 2.2498618
```

# Loops

- We cannot assign new objects in a for loop. Can create a new object set to NULL to represents an empty object.

```
> for(i in 1:n){
+   v9[i] <- v7[i] * 3
+ }
Error in eval(expr, envir, enclos): object 'v9' not found
> # create a empty object
> v9 <- NULL
> v9
NULL
> # build up a vector
> for(i in 1:n){
+   v9[i] <- v7[i] * 3
+ }
> v9
[1] 4.535344 1.169530 -1.863722 -6.644100 3.374793
```

# Loops

```
> # simple projections
> p <- rep(NA, times = 10)
> p[1] <- 100
> r <- 0.05
> for(i in 1:9){
+   p[i+1] <- p[i] * (1 + r)
+ }
> p
[1] 100.0000 105.0000 110.2500 115.7625 121.5506 127.6282 134.0096
[8] 140.7100 147.7455 155.1328
```

# Loops

- Loops can be really useful for repetitive tasks

```
> for(i in 1:10){  
+   # read in different excel sheets  
+   d <- read_excel(path = "myexcelfile.xlsx", sheet = i)  
+   # run regression models using data from excel sheet  
+   m <- lm(formula = y ~ x1 + x2, data = d)  
+   # save regression model coefficients  
+   write_csv(x = m$coefficients, path = paste0("model", i, ".csv"))  
+ }
```

## Exercise 8 (ex18.R)

*# 1. Using the runif() function create an object called r1 based on one random number*

*# 2. Write an if statment to return "Big" if r1 is greater than 0.5*

*# 3. Write an ifelse statment to return "BIG" if r1 is greater than 0.8 and "not big"*

*# 4. Using the runif() function create an object called r2 based on 20 random numbers*

*# 5. Write an ifelse statment to return 1 if r2 is greater than 0.4 and less than 0.6*

*# 6. Create a for loop that prints your name 100 times*

*# 7. In a for loop print the cumulative sum of r2*

*# 8. Create a vector r3 of 20 missing values*

*# 9. In a for loop replace the NA's of r3 with the value of r2 times 2*

*# 10. Create a vector of 10 NA's*

# Functions

- R is a programming language. We can create (program) our own functions using the `function()` function.
  - Wrap up code for repetitive tasks
  - Simplifies complex tasks
  - Creates user friendly access to your R code.
- Functions takes the format

```
function(arguments, ...){  
  actions  
  return(value)  
}
```

- Wrap all sorts of functions, loops, and statements inside of a function to simplify repetitive tasks.
- Specify as many arguments as you like.
- Set argument(s) to `NULL` if we do not want a default value. Forces the user to set.
- The `return()` function is not essential.
  - The actions will take place but nothing is outputted unless we print the object on the last line of the function.



# Functions

```
> sp_rate <- function(p0 = 5, r = 0.05){
+   p <- p0
+   for(i in 1:5){
+     p[i+1] <- p[i] * (1 + r)
+   }
+   return(p)
+ }

> sp_rate()
[1] 5.000000 5.250000 5.512500 5.788125 6.077531 6.381408

> sp_rate(p0 = 20, r = 0.1)
[1] 20.0000 22.0000 24.2000 26.6200 29.2820 32.2102

>

> # add a n argument for the number of projection periods
> sp_rate <- function(p0 = NULL, r = NULL, n = NULL){
+   p <- p0
+   for(i in 1:n){
+     p[i+1] <- p[i] * (1 + r)
+   }
+   return(p)
+ }

> sp_rate(p0 = 20, r = 0.1, n = 20)
[1] 20.00000 22.00000 24.20000 26.62000 29.28200 32.21020 35.43122
[8] 38.97434 42.87178 47.15895 51.87485 57.06233 62.76857 69.04542
[15] 75.94997 83.54496 91.89946 101.08941 111.19835 122.31818 134.55000
```

## Exercise 9 (ex19.R)

```

# 1. Create a function called my_name with inputs
#   a) name
#   b) n
#   that prints your name n times
#   (Hint: Use a for loop)
##### <- function(name, n){
  #####(i in 1:#####){
    #####(name, n)
  }
}

# 2. Run your function with your name and n = 10

# 3. Create a function called rpois_add with inputs
#   a) n
#   b) lambda1
#   c) lambda2
#   that adds together n random numbers from poisson distributions with mean lambda1
rpois_add <- #####(n, #####, lambda2){
  ##### <- rpois(n = #####, lambda = lambda1)
  r2 <- rpois(n = n, ##### = lambda2)
  #####(r1 + #####)
}

# 4. Run your function with n = 20, lambda1 = 5, lambda2 = 8

# 5. Create a function called called_... with inputs

```

# Working Directory

- R works in a particular location on your computer.
- We can change the location to make it easier to
  - Read in data (more on this later)
  - Run R scripts
  - Save data, plots, other outputs (more on this later).
- The `getwd()` function shows the current folder (directory) R is working in.

```
> getwd()
[1] "C:/Users/Guy/Dropbox/SHU2017-3XS371026/slides-md"
```

- The `setwd()` function allows you to set the directory of your R session.
  - Note: R uses / or \\ instead of \ when setting directories:

```
> setwd(dir = "C:/Users/Guy/Dropbox/SHU2017-3XS371026/exercise-solutions")
> # on mac would be
> # setwd(dir = "/Users/Guy/Dropbox/SHU2017-3XS371026/exercise-solutions")
> getwd()
[1] "C:/Users/Guy/Dropbox/SHU2017-3XS371026/exercise-solutions"
> setwd(dir = "C:/Users/Guy/Dropbox/SHU2017-3XS371026/")
> getwd()
[1] "C:/Users/Guy/Dropbox/SHU2017-3XS371026"
```

# RStudio Projects

- With RStudio you can use a project file and never have to worry about `setwd()`
  - File | New Project
    - Select New Directory if you do not have a folder already
    - Select Existing Directory if you do not have a folder with some saved files.
- At the start of your R session use File | Open Project
  - Can also select from the Project dropdown in the top right hand corner.

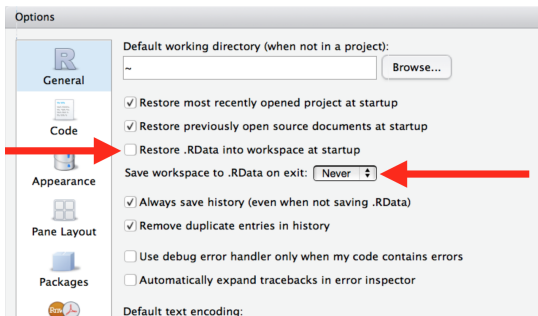
# Sourcing a Script

- The `source()` function runs R code in a file.

```
> # remove all objects
> rm(list = ls())
> # set working directory to solution folder
> setwd(dir = "C:/Users/Guy/Dropbox/SHU2017-3XS371026/exercise-solutions/")
> # run code in ex19.R, could run from anywhere with the full path:
> # source(file= "C:/Users/Guy/Dropbox/SHU2017-3XS371026/exercise-solutions/ex19.R")
> source(file = "ex19.R")
[1] "Guy"
[1] "Guy"
[1] "Guy"
[1] "Guy"
[1] "Guy"
[1] "Guy"
[1] "Guy"
[1] "Guy"
[1] "Guy"
[1] "Guy"
```

# Sourcing a Script

- Using R code scripts (.R files) allows you to save and share your code.
  - Can break big project into multiple scripts.
  - Keeps a tidy workflow
  - Keeps the number of objects to a minimum
  - Quickly identify a particular line of code for revision.
- Instruct RStudio not to preserve your workspace between sessions:
  - Will encourage use of scripts.
  - Easy to run a script to create workspace objects.
  - Difficult to re-create workspace objects without scripts



# Assignment 1 (assign1.R)

```
##
```

```
## Assignment 1
```

```
##
```

```
## Create two questions, with answers below, based on Exercise 1 (ex11.R)
```

```
# 1.
```

```
# 2.
```

```
## Create two questions, with answers below, based on Exercise 2 (ex12.R)
```

```
# 1.
```

```
# 2.
```

```
## Create two questions, with answers below, based on Exercise 3 (ex13.R)
```

```
# 1.
```

```
# 2.
```

```
## Create two questions, with answers below, based on Exercise 4 (ex14.R)
```

```
# 1.
```

# Homework

- For next lesson install the tidyverse set of packages and others using:

```
> install.packages("tidyverse")  
> install.packages("foreign")  
> install.packages("readstata13")  
> install.packages("openxlsx")
```